

Scalable Static Analysis Using Facebook Infer

Dominik Harmim*, Vladimír Marcin**, Ondřej Pavla***



Abstract

Recently, static analysis has become a more popular way of catching bugs early in the modern software. However, while it is quite precise it often fails to scale on bigger codebases. But, the *Facebook Infer*, a static analysis framework, provides a scalable *compositional* and *incremental* solution. We propose to extend Infer with three inter-procedural analyzers: *Looper* (a resource bounds analyser), *L2D2* (a deadlock analyser) and *Atomer* (an atomicity violation analyser). We evaluated our analyzers on set of either artificial examples or official benchmarks and recieved encouraging results. In particular, *L2D2* attained 100 % detection rate and 11 % false positive rate on extensive benchmark of hundreds of functions.

Keywords: Facebook Infer — Static Analysis — Abstract Interpretation — Atomicity Violations — Concurrent Programs — Performance — Worst-case Cost

Supplementary Material: — [Atomer Repository](#) — [Looper Repository](#) 

* xharmi00@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

** xmarci10@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

*** xpavel34@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

1. Introduction

Bugs are an inherent part of a software ever since the inception of the programming discipline. They tend to hide in unexpected places and when they are triggered they can cause a significant damage. In order to catch bugs early in the development process we usually use extensive automated testing or dynamic analysis tools such as profilers. But while these solutions are sufficient in many cases, they fail in many others, some of them unusable for practical projects. We can use as an alternative solution static analysis, however, it has its own shortcomings as well.

Recently, Facebook has proposed its own solution for efficient bug finding and program verification called the *Facebook Infer* — a highly scalable *compositional* and *incremental* framework for creating various inter-procedural analyses. Facebook Infer is still under

development and already provides many various checkers, e.g., for verification of buffer overflow, thread safety or resource leakage, but most of all provides a suitable place for creating new analyses quickly.

However, the current version still misses better support, e.g., for concurrency or performance-based bugs. While it provides a fairly advanced data race and deadlock analyzers, they are limited to Java programs only and fail for C programs, which require more thorough manipulation with locks. Moreover, the only performance-based analyzer focuses only on *worst-case execution time* analysis, which does not provide a reasonable understanding of the programs performance. And while resource bounds analysis and concurrency checkers are not usable for all of the programs, they still can enhance both development process and user experience.

35 We propose to extend the Facebook Infer with
 36 three analyzers: the *Looper*, a resource bounds anal-
 37 yser; the *L2D2*, an lightweight deadlock checker; and
 38 the *Atomer*, an atomicity violation checker. In exper-
 39 imental evaluation we show an encouraging results,
 40 when even our immature implementation could detect
 41 both concurrency property violations and infer precise
 42 bounds for selected benchmarks.

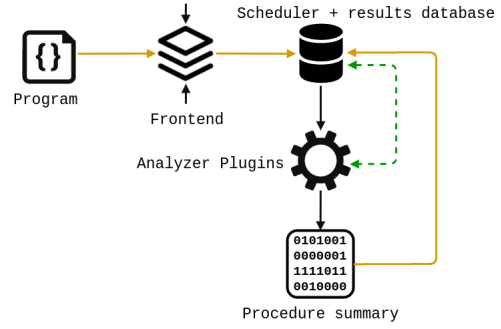


Figure 1. Architecture components

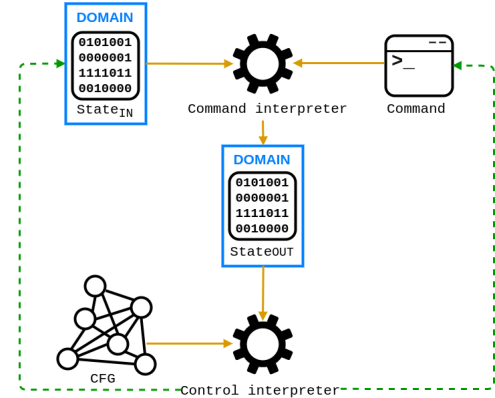


Figure 2. Interpretation process

43 2. Facebook Infer

44 *Facebook Infer* is an open-source static analysis frame-
 45 work which is able to either discover various types of
 46 bugs or verify the correctness of the input programs
 47 both in a *scalable* manner. Infer was originally a
 48 standalone analyzer focused on memory safety viola-
 49 tions which has made its breakthrough thanks to
 50 an influential paper [1]. Since then, it has evolved
 51 into a general abstract interpretation framework that
 52 can be used to quickly develop new kinds of simple
 53 ~~intra-procedural, compositional, incremental or inter-~~
 54 ~~procedural analyses~~ based on the notion of function
 55 summaries. A summary is in theory representation
 56 of function's preconditions, postconditions or effects,
 57 and in practice a custom data structure that allows user
 58 to store arbitrary information after function's analysis.
 59 This way, Infer analyzes each procedure only once
 60 and on demand reuses the summary at multiple call-
 61 sites. Further more, thanks its incrementality Infer
 62 can analyze individual code changes instead of the
 63 whole project, which is more suitable for large and
 64 quickly changing codebases where the conventional
 65 batch analysis is unfeasible.

66 Infer currently supports analysis of multiple lan-
 67 guages including C, C++, Objective-C and Java pro-
 68 grams and provides wide range of analyses each focus-
 69 ing on different types of bugs, such as *Inferbo* (buffer
 70 overruns), *RacerD* (data races) or *Starvation* (concur-
 71 rency starvation and selected types of deadlocks).

72 The architecture of the abstract interpretation (*In-*
 73 *fer-ai*) framework can be divided into three main com-
 74 ponents as depicted in Figure 1: a frontend, an analysis
 75 scheduler and a collective set of analysis plugins.

76 The first component, the front-end, compiles input
 77 programs into the Smallfoot Intermediate Language
 78 (SIL) in form of a Control Flow Graph (CFG). Each
 79 analyzed procedure has its own CFG representation.
 80 The frontend supports multiple languages so one can
 81 write language-independent analyses.

82 The second, the abstract interpreter or *command*
 83 *interpreter*, subsequently interprets SIL instructions
 84 over input abstract states and produces new output state
 85 which are further scheduled for interpretation based

on the CFG. Its simplified workflow is described in
 Figure 2.

The last component, the scheduler determines the
 order of analysis for each procedure based on a *call*
graph and allows Infer to run in a heavily parallelized
 manner as it checks which procedures can be analyzed
 concurrently. Scheduler then stores the results of anal-
 yses in a database for later use in order to ensure the
incremental property of analysis.

Call graph is a directed graph describing call de-
 pendencies between procedures. We can demonstrate
 the analysis order and the incremental property on
 Figure 3.




Figure 3. Call graph

The underlying analyzer starts with leaf proce-
 dures P5 and P6 and then proceeds towards the root
 P_{main} while respecting the dependencies represented
 by the edges. This order ensures that we will always
 have a summary already available when we have to
 abstractly interpret a nested procedure call during our
 analysis.

Each subsequent code change then triggers re-analysis of ~~only~~ directly affected procedures as well as all procedures up the call chain. For example, if we modified the procedure P_3 , Infer will re-analyze only P_3 , P_1 and P_{main} .

3. Worst-case Cost Analyzer

Recently, performance issues ~~has~~ become considerably more widespread in code leading to a poor user experience [2]. This kind of bugs is hard to manifest during the testing and so employing static analysis is ~~nowadays more common~~. Facebook Infer currently provides ~~only~~ the cost checker [3], which implements a *worst-case execution time* complexity analysis (WCET). However, this WCET analysis provides ~~only~~ a numerical bound on number of executions of the program — a bound that is hard to interpret and, ~~most~~ of all, is ~~insufficient~~ for more complex algorithms, e.g., requiring amortized reasoning. Loopus [4] is a powerful resource bounds analyzer, which to the best of our knowledge is the only one that can handle ~~the~~ *amortized complexity analysis* for a broad range of programs. However, Loopus is limited to ~~the~~ intraprocedural analysis only and the tool itself does not scale well. Infer, on the other hand, offering the principles of *compositionality*, can handle even large projects. Hence, recasting the powerful analysis of Loopus within ~~the~~ Infer could enable a more efficient resource bounds analysis usable in today's rapid development.

Cost bounds inferred by Loopus refer to the number of possible *back jumps* to loop headers, which is a useful metric related to *asymptotic time complexity* as it corresponds to the possible number of executions of instructions inside the loop. The bound algorithm relies on a simple abstract program model called *difference constraint program* (DCP) which can be seen in figure 4b.

Listing 1. Snippet demonstrating the need for amortized complexity analysis. Corresponding abstraction in figure 4b. Cost: $3n$

```

142 void foo(int n) {
143     int i = n, j = 0, z = 0;
144      $l_1$ : while (i > 0) {
145         i--; j++;
146     }
147      $l_2$ : while (j > 0 && *) {
148         j--; z++;
149     }
150     int x = z;
151      $l_3$ : while (x > 0)
152         x--;
153 }
```

Each transition τ of a DCP has a *local bound* τ_v

which is a variable v that *locally* limits the number of executions of transition τ as long as some other transitions that might increase the value of v are not executed. For example, the variable j in figure 4b limits the number of consecutive executions of transition τ_2 but not the total number as j might increase on other transitions.

The bound algorithm itself is based on the idea of reasoning about *how often* and *by how much* might the local bound of a transition τ increase which in turn affects the number of executions of τ . There are two main procedures that constitute the algorithm:

1. VB – computes a *variable bound* expression in terms of program parameters which bounds the value variable v .
2. TB – computes a bound on the number of times that a transition τ can be executed. Transitions that are not part of any loop have bound of 1.

The TB procedure is defined in a following way:

$$TB(\tau) = \text{Incr}(\tau_v) + \text{Resets}(\tau_v) \quad (1)$$

The $\text{Incr}(\tau_v)$ procedure implements the idea of reasoning *how often* and *by how much* might the local bound τ_v increase:

$$\sum_{(t,c) \in \mathcal{I}(\tau_v)} TB(t) \times c \quad (2)$$

The $\mathcal{I}(\tau_v)$ is a set of transitions that increase the value of τ_v by c . The $\text{Resets}(\tau_v)$ procedure takes into account the possible resets of local bound τ_v to some arbitrary values which also add to the total amount by which it might increase:

$$\sum_{(t,a,c) \in \mathcal{R}(\tau_v)} TB(t) \times \max(VB(a) + c, 0) \quad (3)$$

The $\mathcal{R}(\tau_v)$ is a set of transitions that reset the value of local bound τ_v to $a + c$ where a is a variable.

The remaining $VB(v)$ procedure is defined as:


$$VB(v) = \text{Incr}(v) + \max_{(t,a,c) \in \mathcal{R}(v)} (VB(a) + c) \quad (4)$$

It picks the maximal value of all possible resets of variable v as an initial value which is subsequently increased by the ~~amount~~ obtained from $\text{Incr}(v)$. Note that the procedure returns v itself if it is a program parameter or a numeric constant.

The complete bound algorithm is thus obtained through the mutual recursion of procedures TB and VB . The main reason why this approach scales so well is *local* reasoning. Loopus does not rely on any

Call	Evaluation and Simplification
$T\mathcal{B}(\tau_5)$	$\rightarrow \text{Incr}([x]) +$ $T\mathcal{B}(\tau_4) \times \max(V\mathcal{B}([z]) + 0, 0)$ $\rightarrow 0 + 1 \times \max([n] + 0, 0) = [n]$
$V\mathcal{B}([z])$	$\rightarrow \text{Incr}([z]) + \max(V\mathcal{B}(0) + 0) = [n]$
$\text{Incr}([z])$	$\rightarrow T\mathcal{B}(\tau_2) \times 1 = [n]$
$T\mathcal{B}(\tau_2)$	$\rightarrow \text{Incr}([j]) + T\mathcal{B}(\tau_0) \times 0$ $\rightarrow [n] + 1 \times 0 = [n]$
$\text{Incr}([j])$	$\rightarrow T\mathcal{B}(\tau_1) \times 1 = [n]$
$T\mathcal{B}(\tau_1)$	$\rightarrow \text{Incr}([i]) + T\mathcal{B}(\tau_0) \times \max([n] + 0, 0)$ $\rightarrow 0 + 1 \times [n] = [n]$

(a) Simplified computation of bound for τ_5 . $\text{Incr}([x])$ and $\text{Incr}([i])$ are 0 as there are no transitions that increase the value of $[x]$ or $[i]$. $T\mathcal{B}(\tau_0)$ and $T\mathcal{B}(\tau_4)$ are 1 as they are not part of any loop.



(b) Abstraction obtained from 1. Each transition is denoted by a set of invariant inequalities.

global program analysis and is able to obtain complex invariants such as $x \leq \max(m1, m2) + 2n$ by means of bound analysis. These invariants are not expressible in common abstract domains such as *octagon* or *polyhedra* which would lead to a less precise result. This approach is also *demand-driven* (4a), which means that it performs *only* necessary recursive calls and does not greedily compute all possible invariants but only the ones that are needed for computation of specified bound. For full flow and path sensitive algorithm and its extension please refer to [4].

The table 4a presents simplified computation of transition bound of τ_5 from DCP 4b which was obtained through abstraction algorithm from the code snippet 1. This code snippet demonstrates the need for amortized complexity analysis as the worst-case cost of the l_2 loop can indeed be n . However, the amortized cost is 1 because the total number of iterations (total cost) is also equal to n due to the local bound j which is bounded by n . Loopus is thus able to obtain bound of n instead of n^2 for the inner loop l_2 unlike many other tools that cannot reason about amortized complexity. Another challenging problem is the computation of bound for the loop l_3 . It is easy to infer z as the bound but the real challenge lies in expressing the bound in terms of program parameters. Thus, the real task is to obtain an invariant of form $z \leq \text{expr}(n)$ where $\text{expr}(n)$ denotes an expression over program parameters, n in this case. Loopus is able to obtain the invariant $z \leq n$ simply with the $V\mathcal{B}$ procedure and consequently infer the bound n for the loop l_3 .

The table 1 presents results which we were able to achieve with our current implementation on few artificial examples. We compared the results of *Looper*

(Loopus in Infer) with the *Cost* analyzer mentioned in the introduction of this section. Please note that the real cost of examples #4 and #6 is in fact $n \times \max(n - 1, 0) + n$ and $3n + \max(m1, m2)$. Displayed cost of these examples is actually the worst-case asymptotic complexity instead of cost.

	Bound	Inferred bound		Time [s]	
		Looper	Cost	Looper	Cost
#1	n	$2n$	—	0.3	—
#2	$2n$	$2n$	—	0.5	—
#3	$4n$	$5n$	—	0.8	—
#4	n^{2*}	n^2	—	0.6	—
#5	$2n$	$2n$	—	0.3	—
#6	n^*	n	—	0.6	—
#7	$2n$	$2n$	—	0.4	—
#8	$2n$	$2n$	—	0.7	—

Table 1. Experimental evaluation on selected examples used for evaluation of Loopus [4]. Benchmarks are publicly available at [bitbucket](https://bitbucket.org).

4. Deadlock Analyzer

According to [5] deadlock is perhaps the most common concurrency error that might occur in almost all parallel programming paradigms including both shared-memory and distributed memory. To detect deadlock during testing is very hard due to many possible interleavings between threads that's the reason why many of detectors were created, but most of them are quite heavyweight and do not scale well. However, there are a few that meet the scalability condition, like *starvation* analyzer implemented in Facebook Infer. The problem of this analyzer is that it uses heuristic

on the class room the access path of the lock so it doesn't handle a pure C lock. Also worth mentioning is the RacerX analyzer [6], which is based on counting so-called *locksets* i.e. sets of locks currently held. RacerX uses interprocedural, flow-sensitive and context-sensitive analysis. What means that each function needs to be reanalysed in a new context. Hence, we decide to adapt lockset analysis from RacerX to follow principles of Facebook Infer and by that create context-insensitive analysis which will be faster and more scalable. So we present Low Level Deadlock Detector (L2D2), the principle of which will be illustrated with the example in Listing 2.

L2D2 works by first computing a summary for each function by looking for lock and unlock events. Example of lock and unlock is illustrated in Listing 2 at lines 22 and 27. If user function `call` appears in the analyzed code during analysis, like at line 26 of our example, the analyzer is provided with a summary of the function if available or the function is analyzed on demand. The summary is then applied to an abstract state at a call site. So in our example summary of `foo` will be applied to the abstract state of `thread1`.

Listing 2. Simple example capturing a deadlock between two global locks in C language using POSIX threads execution model

```

269 16 void foo() {
270 17     pthread_mutex_lock(&lock2);
271 18 }
272 21 void *thread1(...) {
273 22     pthread_mutex_lock(&lock1);
274 23     :
275 26     foo();
276 27     pthread_mutex_unlock(&lock1);
277 28 }
278 29 void *thread2(...) {
279 30     pthread_mutex_lock(&lock2);
280 31     :
281 36     pthread_mutex_lock(&lock1);
282 37 }

```

Next L2D2 looks through all the summaries of analyzed program and checks whether a potential deadlock can occur by computing transitive closure of relation consisting of all dependencies (see Listing 3) and checking if any lock depends on itself. The summaries for functions from the above example record information about the state of locks `lock1` and `lock2` as follows:

Listing 3. Summaries of the functions in Listing 2

```

291 foo()
292     PRECONDITION: { unlocked={lock2} }
293     POSTCONDITION: { lockset={lock2} }
294 thread1(...)
295     PRECONDITION: { unlocked={lock1, lock2} }

```

```

296 POSTCONDITION: {
297     lockset={lock1, lock2},
298     dependencies={lock1->lock2}
299 }
300 thread2(...)
301     PRECONDITION: { unlocked={lock1, lock2} }
302     POSTCONDITION: {
303         lockset={lock1, lock2},
304         dependencies={lock2->lock1}
305     }

```

If we run L2D2 on code from our example it will report a possible deadlock between two threads due to cyclic dependency between `lock1->lock2` and `lock2->lock1` that arises if thread 1 holds `lock1` and waits on `lock2` and thread 2 holds `lock2` and waits on `lock1`.

4.1 Computing procedure summaries

In this subsection, we describe structure of the summary and process of computing it. To detect potential deadlock we need to record information that will allow us to answer these questions:

- (1) What is the state of locks used in the analyzed program?
- (2) Could cyclic dependency on pending threads occur?

To answer question (1), we have defined sets *lockset* and *unlockset*, which contains currently locked and unlocked locks respectively. We have also added sets *locked* and *unlocked* that serve as a precondition for a given function and contain locks that should be locked/unlocked before calling this function. Semantic of these sets is as follows:

```

328 semantics of lockset:
329     lock(l) → lockset = lockset ∪ {l}
330     unlock(l) → lockset = lockset - {l}
331 semantics of unlockset:
332     lock(l) → unlockset = unlockset - {l}
333     unlock(l) → unlockset = unlockset ∪ {l}
334 semantics of locked:
335     if(lock(l) is first operation in f)
336         unlockedf = unlockedf ∪ {l}
337 semantics of unlocked:
338     if(unlock(l) is first operation in f)
339         lockedf = lockedf ∪ {l}

```

The summary also contains a set of *one-level dependencies* by using which we can answer (2)nd question. Extraction of these dependencies is called on every lock acquisition and iterates over every lock in the current *lockset*, emitting the ordering constraint produced by the current acquisition. For example, if `lock2` is in the current *lockset* and `lock1` has just been acquired, the dependency `lock2->lock1` will be emitted, as we can see in Listing 2 in function `thread2`.

350 The most difficult part of dependencies extraction
 351 is elimination of false ones caused by invalid locksets.
 352 The main reasons for errors in locksets include the
 353 number of conditionals, function calls and degree of
 354 aliasing involved.

355 Applying procedure summaries. As we mentioned
 356 at the beginning of this section, if a function call ap-
 357 pears in an analyzed code, we have to apply a summary
 358 of the function to an abstract state at a callsite. Given
 359 callee g , its lockset L_g , unlockset U_g and caller f , its
 360 lockset L_f , unlockset U_f and dependencies D_f , we:

- 361 (1) Update the summary of g by replacing formal
 362 parameters with actual ones in case that locks
 363 were passed to g as parameters. In the example
 364 below, you can notice that in the summary of g
 365 will be `lock4` replaced with `lock2`.
- 366 (2) Update the precondition of f :
 367 $if(\exists l : l \in unlocked_g \wedge l \notin unlockset_f)$
 368 add lock l to $unlocked_f$
 369 $if(\exists l : l \in locked_g \wedge l \notin lockset_f)$
 370 add lock l to $locked_f$
- 371 (3) Update L_f : $L_f = (L_f \setminus U_g) \cup L_g$
- 372 (4) Update U_f : $U_f = (U_f \setminus L_g) \cup U_g$
- 373 (5) Update D_f by adding new dependencies for all
 374 locks in the L_g with locks which were locked in
 375 g . But what if all the locks which were acquired
 376 in g were also released there, as we can see in
 377 the example below.

```

378 void f() {
379     pthread_mutex_lock(&lock2);
380     g(&lock2);
381 }
382 void g(pthread_mutex_t *lock4) {
383     pthread_mutex_lock(&lock3);
384     pthread_mutex_unlock(lock4);
385     pthread_mutex_lock(&lock1);
386     :
387     pthread_mutex_unlock(&lock1);
388     pthread_mutex_unlock(&lock3);
389 }
```

390 In that case, L_g will be empty and we have no
 391 information about these locks. So we had to add
 392 a new set to the summary which semantics is
 393 similar to the semantics of lockset except that
 394 unlock statement does not remove a lock from it.
 395 In our example, this set would contain `lock3`
 396 and `lock1` but there is still one problem left.
 397 What if the lock from the current lockset was
 398 unlocked in the callee before we locked another
 399 lock there? Then we will emit the wrong depen-
 400 dency `lock2→lock1`. In order to avoid this,
 401 we create `unlock→lock` type dependencies
 402 in summary, that can be used to safely determine

the order of operations in the callee. So this en-
 403 sures that the only newly created correct depen-
 404 dency in our example will be `lock2→lock3`. 405

4.2 Reporting deadlocks 406

407 For deadlock detection, we use algorithm that iterates
 408 through all the summaries and computes the transitive
 409 closure of all dependencies. It records the cyclic lock
 410 dependency and displays the results to the user for
 411 inspection. Each deadlock is normally reported twice,
 412 at each trace starting point. So in our example in
 413 Listing 2, will be the deadlock reported for the first
 414 time in function `thread1` and for the second time in
 415 function `thread2`.

4.3 Experimental evaluation 416

417 We performed our experiments by using 1002 concu-
 418 rent C programs, that contain locks from the Debian
 419 GNU/Linux distribution. All benchmarks are avail-
 420 able online at [gitlab](#). These programs were used for
 421 experimental evaluation of Daniel Kroening's static
 422 deadlock analyser [7] implemented in the CPROVER
 423 framework.

424 This benchmark set consists of 11.4 MLOC. Of
 425 all the programs, 994 are assumed to be deadlock-free
 426 and 8 of them have proved deadlock. Our experiments
 427 were run on a CORE i7-7700HQ at 2.80 GHz running
 428 Ubuntu 18.04 with 64-bit binaries with comparison
 429 to the CPROVER experiments which were run on a
 430 Xeon X5667 at 3 GHz running Fedora 20 with 64-
 431 bit binaries. In case of CPROVER were memory and
 432 CPU time restricted to 24GB and 1800 seconds per
 433 benchmark.

434 Results. Our analyzer as same as CPROVER cor-
 435 rectly report all 8 potential deadlocks in benchmarks
 436 with known issues. Comparison of results for deadlock-
 free programs you can see in Table 2.

	proved	alarms	t/o	m/o	errors
CPROVER	292	114	453	135	0
L2D2	810	104	0	0	80

Table 2. Results for the programs without deadlock
 (t/o – timed out, m/o – out of memory)

437 As you can see L2D2 reported false alarms for
 438 104 deadlock-free benchmarks what is 10 less than
 439 CPROVER. A much larger difference can be seen in
 440 cases where it was proved that there was no deadlock.
 441 The difference here is up to 518 examples in favor of
 442 our analyzer. In case of L2D2 you can see 80 com-
 443 pilation errors that were caused by syntax that Infer
 444 does not support. The biggest difference between our
 445 analyzer and CPROVER is runtime. While our ana-
 446

lyzer needed approximately 2 hours to perform the experiments, CPROVER needed about 300 hours.

There is still space for improving our analysis by reduction of false alarms. The main reason for such alarms is false dependencies. Reasons for their existence we mentioned in subsection 4.1 (4th paragraph). So eliminating false positives consists of techniques to eliminate false dependencies. Some techniques have already been implemented but we are still working on others.

5. Atomicity Violations Analyzer

In concurrent programs, there are often *atomicity requirements* for an execution of specific sequences of instructions. Violating these requirements may cause many kinds of problems, such as a unexpected behaviour, exceptions, segmentation faults or other failures. *Atomicity violations* are usually not verified by compilers, unlike syntactic or some sorts of semantic rules. Atomicity requirements, in most cases, are not event documented. It means that typically only programmers must take care of following these requirements. In general, it is very difficult to avoid errors in *atomicity-related programs*, especially in large projects, and even harder and time-consuming is then finding and fixing these errors.

In this section of this paper, there is described a proposal and an implementation of an *static analyzer for finding atomicity violations*.

5.1 Contracts for Concurrency

The proposal of a solution is based on the concept of *contracts for concurrency* described in [?]. These contracts allow to define *sequences of functions* that are required to be *executed atomically*. The proposed analyzer itself (**Atomer**) is able to produce mentioned contracts, and then verify whether the contracts are fulfilled.

In [?], a *basic contract* is formally defined as follows. Let $\Sigma_{\mathbb{M}}$ be a set of all function names of a software module. A *contract* is a set \mathbb{R} of *clauses* where each clause $\varrho \in \mathbb{R}$ is a regular expression over $\Sigma_{\mathbb{M}}$. A *contract violation* occurs if any of the sequences represented by the contract clauses is interleaved with an execution of functions from $\Sigma_{\mathbb{M}}$.

Consider an implementation of a function that replaces item *a* in an array by item *b*, as illustrates Listing 4. The contract for this specific scenario contains clause ϱ_1 , which is defined and follows:

Clause ϱ_1 specifies that the execution of `index_of` followed by execution of `set` should be atomic. The index of an item in an array is acquired, and then the index is used to modify the array. Without atomicity, a concurrent modification of the array may change a position of the item. The acquired index then may be invalid when `set` is executed.

Listing 4. Example of a contract violation

```
void replace(int *array, int a, int b) {  
    int i = index_of(array, a);  
    if (i >= 0) set(array, i, b);  
}
```

In paper [?], there is described a proposal and an implementation for a static validation which is based on *grammars* and *parsing trees*. Within paper [?] was implemented a stand-alone prototype tool¹ for analysing programs written in Java language, which obtained promising experimental results. However, we decided to propose and implement the analysis quite different way, see 5.2 and 5.3. Moreover, we decided to implement this solution in the *Facebook Infer*, i.e., widely used, active and a open source tool. Therefore the analysis should be faster and more scalable thanks to the way how the Facebook Infer works, as it was described in section 2. The implementation is aimed for programs written in C/C++ languages using *POSIX Threads (Pthreads)* locks for a *synchronization of concurrent threads*. We are also focusing to reduce false positive errors.

In the Facebook Infer, there is already implemented an analysis called *Lock Consistency Violation*, see ², which is part of the *RacerD* [?]. That analysis finds atomicity violations for writes/reads single variables that are required to be executed atomically. *Atomer* is more general because it finds *atomicity violations for sequences of functions* that are required to be executed atomically, i.e., it checks whether contracts for concurrency are fulfilled.

The proposed solution is divided into two parts (*phases of analysis*):

- Phase 1** A detection of *atomic sequences* 5.2
- Phase 2** A detection of *atomicity violations* 5.3.

5.2 Detection of Atomic Sequences

[[TODO]]

5.3 Detection of Atomicity Violations

[[TODO]]

¹<https://github.com/trxsys/gluon>

²https://fbinfer.com/docs/checkers-bug-types.html#LOCK_CONSISTENCY_VIOLATION

(ϱ_1) `index_of set`

6. Conclusions

In this paper, we presented three new analyzers which we implemented in the *Facebook Infer* tool alongside the existing ones. The *Looper* resource bounds analyzer was able to infer the precise bound in 6 out of 8 of selected examples used for evaluation of the original *Loopus* tool. The remaining two bounds differed only in the constant factor. The *L2D2* analyzer focusing on deadlock detection in C programs was evaluated on Daniel Kroening's benchmark with 100 % success rate in detection of potential deadlocks and roughly 11 % false positives rate. It also proved the scalability of the approach as it managed to finish the benchmark in less than 1 % of the time needed by the Kroening's CPROVER tool. The first phase of the *Atomer* – the atomicity violations analyzer, a detection of sequences of functions that should be executed atomically, was successfully verified on a set of sample programs created for this purpose. The second phase, a detection of atomicity violations, will be finalized and tested within a bachelor's thesis.

Our analyzers have potential for further extending and improving the accuracy of their results. So our further work will focus mainly on increasing the accuracy of our methods, and testing them on real-world programs. Furthermore we would like to merge our implementations to a master branch of the *Facebook Infer repository*³.

Acknowledgements

We would like to thank our supervisors and colleagues from VeriFIT team – Tomáš Fiedor, Adam Rogalewicz and Tomáš Vojnar. We would also like to thank for the support received from the H2020 ECSEL project Aquas.

References

- [1] Cristiano Calcagno, Dino Distefano, Peter O'Hearn, and Hongseok Yang. *Compositional Shape Analysis by means of Bi-Abduction*. In *Proc. of POPL'09*, pages 289–300. 2009.
- [2] Guoliang Jin, Linhai Song, Xiaoming Shi, Joel Scherpelz, and Shan Lu. *Understanding and Detecting Real-World Performance Bugs*. In *Proc. of PLDI'18*, pages 77–88. 2012.
- [3] Stefan Bygde. *Static WCET analysis based on abstract interpretation and counting of elements*. PhD thesis, Mälardalen University, 2010.

- [4] Moritz Sinn. *Automated Complexity Analysis for Imperative Programs*. PhD thesis, Vienna University of Technology, 2016.
- [5] A.H. Dogru and V. Bicar. *Modern Software Engineering Concepts and Practices: Advanced Approaches*. Premier reference source. Information Science Reference, 2011.
- [6] Dawson R. Engler and Ken Ashcraft. *Racerx: Effective, static detection of race conditions and deadlocks*. In *Operating Systems Review – SIGOPS*, volume 37, pages 237–252, January 2003. DOI: 10.1145/1165389.945468.
- [7] D. Kroening, D. Poetzl, P. Schrammel, and B. Wachter. *Sound static deadlock analysis for c/pthreads*. In *2016 31st IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 379–390, Sep. 2016.
- [8] Ricardo J. Dias, Carla Ferreira, Jan Fiedor, João M. Lourenço, Aleš Smrčka, Diogo G. Sousa, and Tomáš Vojnar. *Verifying concurrent programs using contracts*. In *2017 IEEE International Conference on Software Testing, Verification and Validation (ICST)*, pages 196–206. IEEE, March 2017. DOI: 10.1109/ICST.2017.25. ISBN: 9781509060313.

³<https://github.com/facebook/infer>