

Scalable Static Analysis Using Facebook Infer

Dominik Harmim*, Vladimír Marcin**, Ondřej Pavla***



Abstract

What is the problem? What is the topic?, the aim of this paper? Lorem ipsum dolor sit amet, consectetur adipiscing elit. Fusce ullamcorper suscipit euismod. Mauris sed lectus non massa molestie congue. In hac habitasse platea dictumst. How is the problem solved, the aim achieved (methodology)? Lorem ipsum dolor sit amet, consectetur adipiscing elit. Fusce ullamcorper suscipit euismod. Mauris sed lectus non massa molestie congue. In hac habitasse platea dictumst. Curabitur massa neque, commodo posuere fringilla ut, cursus at dui. Nulla quis purus a justo pellentesque. What are the specific results? How well is the problem solved? Lorem ipsum dolor sit amet, consectetur adipiscing elit. Fusce ullamcorper suscipit euismod. Mauris sed lectus non massa molestie congue. In hac habitasse platea dictumst. So what? How useful is this to Science and to the reader? Lorem ipsum dolor sit amet, consectetur adipiscing elit. Fusce ullamcorper suscipit euismod.

Keywords: Facebook Infer — Static Analysis — Abstract Interpretation — Atomicity Violations — Concurrent Programs

Supplementary Material: [Facebook Infer](#) — [Facebook Infer Repository](#) — [Atomicity Violations Analyzer Repository](#)

* xharmi00@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

** xmarci10@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

*** xpavel34@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

1. Introduction

Bugs are inherent part of software since the inception of programming discipline. Current solutions to this problem mostly include extensive automated testing and dynamic analysis tools such as profilers. These solutions are sufficient in many cases but mostly fail in others or they are outright not usable such as in early development cycles without a running prototype. Static analysis can provide complementing alternative which however historically had shortcoming in scalability department. *Facebook Infer* provides a solution to this problem with its highly scalable *compositional* and *incremental* inter-procedural analyses.

However, the current version of Infer still lacks in concurrency and performance areas. It provides fairly advanced data race analysis and relatively new deadlock analysis which unfortunately does not work well on C programs. The only performance oriented analyzer focuses on *worst-case execution time* analysis that has certain drawbacks which we will cover later.

[Motivation] What is the raison d'être of your project? Why should anyone care? No general meaningless claims. Make bulletproof arguments for the importance of your work. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer sit amet neque vel mi sodales interdum nec a mi. Aliquam eget turpis

27 venenatis, tincidunt purus eget, euismod neque. Nulla
28 et porta tortor, id lobortis turpis. Sed scelerisque sem
29 eget ante interdum, vel volutpat arcu volutpat.

30 **[Problem definition]** What exactly are you solv-
31 ing? What is the core and what is a bonus? What
32 parameters should a proper solution of the problem
33 have? Define the problem precisely and state how its
34 solution should be evaluated. Lorem ipsum dolor sit
35 amet, consectetur adipiscing elit. Pellentesque non
36 arcu quis nunc efficitur vestibulum. Integer gravida
37 neque suscipit diam porta aliquet. Maecenas porttitor
38 libero ut turpis porttitor, auctor porta ligula rhoncus.
39 Etiam a turpis blandit, eleifend dolor eget, egestas
40 ligula. Nullam sollicitudin pulvinar mi sit amet in-
41 terdum. Etiam in ultrices ante. Suspendisse potenti.
42 Duis vel nisi eget tellus volutpat tempor. Etiam laoreet
43 magna elit, et sollicitudin lectus tempor sit. Maecenas
44 porttitor libero ut turpis porttitor, auctor porta ligula
45 rhoncus. Etiam a turpis blandit, eleifend dolor eget,
46 egestas ligula.

47 **[Existing solutions]** Discuss existing solutions, be
48 fair in identifying their strengths and weaknesses. Cite
49 important works from the field of your topic. Try to
50 define well what is the *state of the art*. You can in-
51 clude a Section 2 titled “Background” or “Previous
52 Works” and have the details there and make this para-
53 graph short. Or, you can enlarge this paragraph to a
54 whole page. In many scientific papers, *this* is the most
55 valuable part if it is written properly. Lorem ipsum
56 dolor sit amet, consectetur adipiscing elit. Praesent
57 congue enim eu eros dictum sagittis. Aliquam ligula
58 arcu, gravida at augue et, aliquet condimentum nulla.
59 Morbi a lectus arcu. Nam ac commodo nisi, a accum-
60 san nunc. Nam sed ante vel nulla elementum lobortis.
61 Aliquam sed laoreet risus. Etiam ipsum odio, gravida
62 eget sapien dictum, eleifend aliquet ex. Duis dapibus
63 vitae enim vitae bibendum. Phasellus eget pulvinar
64 massa. Mauris ornare urna. Maecenas porttitor libero
65 ut turpis porttitor, auctor porta ligula rhoncus. Etiam a
66 turpis blandit, eleifend dolor eget, egestas ligula. Nul-
67 lam sollicitudin pulvinar mi sit amet interdum. Etiam
68 in ultrices ante. Suspendisse potenti. Duis vel nisi eget
69 tellus volutpat tempor. Suspendisse potenti. Duis vel
70 nisi eget tellus volutpat tempor.

71 **[Our solution]** Make a quick outline of your ap-
72 proach – pitch your solution. The solution will be
73 described in detail later, but give the reader a very
74 quick overview now. Lorem ipsum dolor sit amet, con-
75 sectetur adipiscing elit. Morbi laoreet risus a egestas
76 imperdiet. Ut egestas nibh non fermentum vestibulum.
77 Nullam quis eleifend ex, sed maximus nisl. Mauris
78 maximus non dolor id tristique. Nunc pulvinar congue

gravida. Nullam lobortis viverra leo sed commodo. 79
Nulla in elit congue, ullamcorper metus non, eleifend 80
risus. Vivamus porttitor, ex nec porttitor pretium, 81
libero turpis ultrices dui, eu efficitur ante ipsum vel 82
justo. Vivamus nec nulla nisi. Aenean quis mauris 83
vitae metus gravida congue. 84

[Contributions] Sell your solution. Pinpoint your 85
achievements. Be fair and objective. Lorem ipsum 86
dolor sit amet, consectetur adipiscing elit. Integer sit 87
amet neque vel mi sodales interdum nec a mi. Aliquam 88
eget turpis venenatis, tincidunt purus eget, euismod 89
neque. Nulla et porta tortor, id lobortis turpis. Sed 90
scelerisque sem eget ante interdum, vel volutpat arcu 91
volutpat. Aliquam cursus, dolor a luctus. 92

2. Facebook Infer 93

Facebook Infer is an open-source static analysis tool 94
which is able to discover inter-procedural bugs in a 95
scalable manner. Infer was originally a standalone 96
analyzer focused on memory safety violations which 97
has made its breakthrough thanks to the influential pa- 98
per [1]. Since then it has evolved into a general abstract 99
interpretation framework that can be used to quickly 100
develop new kinds of simple intra-procedural or *com-* 101
positional and *incremental* inter-procedural analyses 102
based on *summaries*. Summary is a custom data struc- 103
ture that stores usually parametric information about 104
the analyzed procedure. It allows Infer to analyze each 105
procedure only once and then reuse the acquired info 106
at multiple callsites. The incremental property refers 107
to the ability to analyze individual code changes which 108
makes Infer suitable for large and quickly changing 109
codebases where the conventional batch analysis is 110
unfeasible. 111

It currently supports analysis of C, C++, Objective- 112
C and Java programs and provides wide range of anal- 113
yses each focusing on different bug types. List of 114
more matured analyses includes for example *Inferbo* 115
(buffer overruns), *RacerD* (data races) or *Starvation* 116
(concurrency starvation and some types of deadlocks). 117

The simplified architecture of the abstract interpre- 118
tation (AI) framework consisting of three main com- 119
ponents can be seen in figure 1. 120

The front-end is responsible for compilation from 121
a source language into the Smallfoot Intermediate Lan- 122
guage (SIL) used by AI and generation of Control Flow 123
Graph (CFG) for each analyzed procedure. Thanks to 124
the front-end module, it is possible to write language- 125
independent analyses. 126

The abstract interpreter is responsible for the anal- 127
ysis of individual procedures. It takes a CFG and a 128
module implementing the effect of each SIL instruc- 129

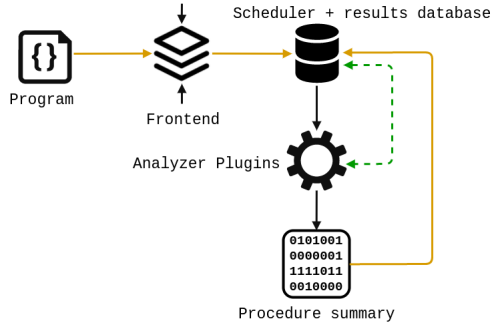


Figure 1. Architecture components

tion and produces a summary. The interpretation process is described in figure 2. The *command interpreter*

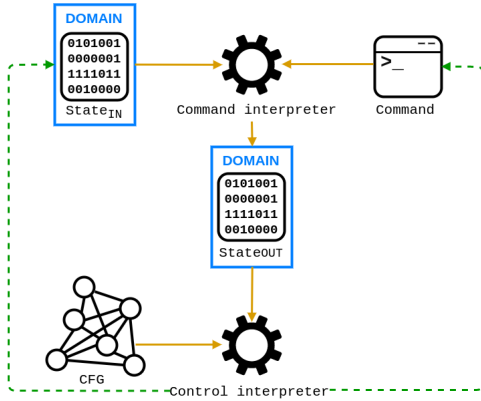


Figure 2. Interpretation process

interprets SIL instruction over the input state, produces new output state and sends it to the *control interpreter* which provides next input state and instruction based on a CFG.

The scheduler determines the order of analysis for each procedure based on a *call graph* and allows Infer to run in a heavily parallelized manner as it checks which procedures can be analyzed concurrently. Scheduler then stores the results in a database for later use in order to ensure the *incremental* property of analysis.

Call graph is a directed graph describing call dependencies between procedures. We can demonstrate the analysis order and the incremental property on figure 3.

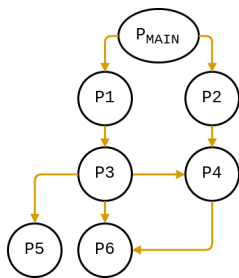


Figure 3. Call graph

The initial run would start with leaf procedures P5 and P6 and then would proceed towards the root P_{main} while respecting the dependencies represented

by edges. This order ensures that we will always have a summary already available when we encounter a procedure call during the analysis.

Each subsequent code change would then trigger re-analysis of directly affected procedures and also all procedures up the call chain. For example, if we modified the procedure P3, Infer would have to re-analyze only the P3, P1 and P_{main}.

3. Worst-case cost analyzer

Recently, performance issues has become considerably more widespread in code leading to a poor user experience [2]. This kind of bugs is hard to manifest during the testing and so employing static analysis is nowadays more common. Facebook Infer currently provides only the *cost* checker [3], which implements a *worst-case execution time* complexity analysis (WCET). However, this WCET analysis provides only a numerical bound on number of executions of the program — a bound that is hard to interpret and, most of all, is insufficient for more complex algorithms, e.g., requiring amortized reasoning. Loopus [4] is a powerful resource bounds analyzer, which to the best of our knowledge is the only one that can handle the *amortized complexity analysis* for a broad range of programs. However, Loopus is limited to the intraprocedural analysis only and the tool itself does not scale well. Infer, on the other hand, offering the principles of *compositionality*, can handle even large projects. Hence, recasting the powerful analysis of Loopus within the Infer could enable a more efficient resource bounds analysis usable in today's rapid development.

Cost bounds inferred by Loopus refer to the number of possible *back jumps* to loop headers which is a useful metric related to *asymptotic time complexity* as it corresponds to the possible number of executions of instructions inside the loop. The bound algorithm relies on a simple abstract program model called *difference constraint program* (DCP) which can be seen in figure 4b.

Listing 1. Snippet demonstrating the need for amortized complexity analysis. Corresponding abstraction in figure 4b. Cost: $3n$

```
void foo(int n) {
    int i = n, j = 0, z = 0;
    l1: while (i > 0) {
        i--; j++;
        l2: while (j > 0 && *) {
            j--; z++;
        }
    }
    int x = z;
    l3: while (x > 0)
```

```

198     x--;
199 }

```

Each transition τ of a DCP has a *local bound* τ_v which is a variable v that *locally* limits the number of executions of transition τ as long as some other transitions that might increase the value of v are not executed. For example, the variable j in figure 4b limits the number of consecutive executions of transition τ_2 but not the total number as j might increase on other transitions.

The bound algorithm itself is based on the idea of reasoning about *how often* and *by how much* might the local bound of a transition τ increase which in turn affects the number of executions of τ . There are two main procedures that constitute the algorithm:

1. VB – computes a *variable bound* expression in terms of program parameters which bounds the value variable v .
2. TB – computes a bound on the number of times that a transition τ can be executed. Transitions that are not part of any loop have bound of 1.

The TB procedure is defined in a following way:

$$TB(\tau) = \text{Incr}(\tau_v) + \text{Resets}(\tau_v) \quad (1)$$

The $\text{Incr}(\tau_v)$ procedure implements the idea of reasoning *how often* and *by how much* might the local bound τ_v increase:

$$\sum_{(t,c) \in \mathcal{I}(\tau_v)} TB(t) \times c \quad (2)$$

The $\mathcal{I}(\tau_v)$ is a set of transitions that increase the value of τ_v by c . The $\text{Resets}(\tau_v)$ procedure takes into account the possible resets of local bound τ_v to some arbitrary values which also add to the total amount by which it might increase:

$$\sum_{(t,a,c) \in \mathcal{R}(\tau_v)} TB(t) \times \max(VB(a) + c, 0) \quad (3)$$

The $\mathcal{R}(\tau_v)$ is a set of transitions that reset the value of local bound τ_v to $a + c$ where a is a variable.

The remaining $VB(v)$ procedure is defined as:

$$VB(v) = \text{Incr}(v) + \max_{(t,a,c) \in \mathcal{R}(v)} (VB(a) + c) \quad (4)$$

It picks the maximal value of all possible resets of variable v as an initial value which is subsequently increased by the amount obtained from $\text{Incr}(v)$. Note that the procedure returns v itself if it is a program parameter or a numeric constant.

The complete bound algorithm is thus obtained through the mutual recursion of procedures TB and

VB . The main reason why this approach scales so well is *local* reasoning. Loopus does not rely on any global program analysis and is able to obtain complex invariants such as $x \leq \max(m1, m2) + 2n$ by means of bound analysis. These invariants are not expressible in common abstract domains such as *octagon* or *polyhedra* which would lead to a less precise result. This approach is also *demand-driven* (4a) which means that it performs only necessary recursive calls and does not greedily compute all possible invariants but only the ones that are needed for computation of specified bound. For full *flow* and *path sensitive* algorithm and its extension please refer to [4]

The table 4a presents simplified computation of transition bound of τ_5 from DCP 4b which was obtained through abstraction algorithm from the code snippet 1. This code snippet demonstrates the need for amortized complexity analysis as the worst-case cost of the l_2 loop can indeed be n . However, the amortized cost is 1 because the total number of iterations (total cost) is also equal to n due to the local bound j which is bounded by n . Loopus is thus able to obtain bound of n instead of n^2 for the inner loop l_2 unlike many other tools that cannot reason about amortized complexity. Another challenging problem is the computation of bound for the loop l_3 . It is easy to infer z as the bound but the real challenge lies in expressing the bound in terms of program parameters. Thus, the real task is to obtain an invariant of form $z \leq \text{expr}(n)$ where $\text{expr}(n)$ denotes an expression over program parameters, n in this case. Loopus is able to obtain the invariant $z \leq n$ simply with the VB procedure and consequently infer the bound n for the loop l_3 .

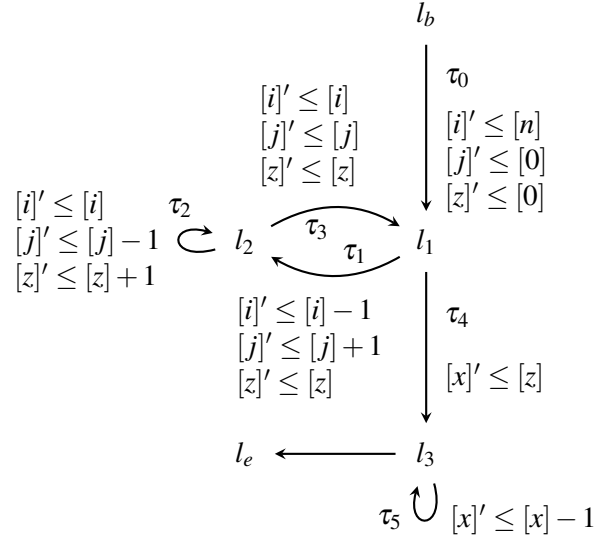
The table 1 presents results which we were able to achieve with our current implementation on few artificial examples. We compared the results of *Looper* (Loopus in Infer) with the *Cost* analyzer mentioned in the introduction of this section. Please note that the real cost of examples #4 and #6 is in fact $n \times \max(n - 1, 0) + n$ and $3n + \max(m1, m2)$. Displayed cost of these examples is actually the worst-case asymptotic complexity instead of cost.

4. Deadlock analyzer

According to [5] deadlock is perhaps the most common concurrency error that might occur in almost all parallel programming paradigms including both shared-memory and distributed memory. To detect deadlock during testing is very hard due to many possible interleavings between threads. That's the reason why many of detectors were created, but most of them are quite heavyweight and do not scale well. However,

Call	Evaluation and Simplification
$T\mathcal{B}(\tau_5)$	$\rightarrow \text{Incr}([x]) +$ $T\mathcal{B}(\tau_4) \times \max(V\mathcal{B}([z]) + 0, 0)$ $\rightarrow 0 + 1 \times \max([n] + 0, 0) = [n]$
$V\mathcal{B}([z])$	$\rightarrow \text{Incr}([z]) + \max(V\mathcal{B}(0) + 0) = [n]$
$\text{Incr}([z])$	$\rightarrow T\mathcal{B}(\tau_2) \times 1 = [n]$
$T\mathcal{B}(\tau_2)$	$\rightarrow \text{Incr}([j]) + T\mathcal{B}(\tau_0) \times 0$ $\rightarrow [n] + 1 \times 0 = [n]$
$\text{Incr}([j])$	$\rightarrow T\mathcal{B}(\tau_1) \times 1 = [n]$
$T\mathcal{B}(\tau_1)$	$\rightarrow \text{Incr}([i]) + T\mathcal{B}(\tau_0) \times \max([n] + 0, 0)$ $\rightarrow 0 + 1 \times [n] = [n]$

(a) Simplified computation of bound for τ_5 . $\text{Incr}([x])$ and $\text{Incr}([i])$ are 0 as there are no transitions that increase the value of $[x]$ or $[i]$. $T\mathcal{B}(\tau_0)$ and $T\mathcal{B}(\tau_4)$ are 1 as they are not part of any loop.



(b) Abstraction obtained from 1. Each transition is denoted by a set of invariant inequalities.

	Bound	Inferred bound		Time [s]	
		Looper	Cost	Looper	Cost
#1	n	$2n$	—	0.3	—
#2	$2n$	$2n$	—	0.5	—
#3	$4n$	$5n$	—	0.8	—
#4	n^{2*}	n^2	—	0.6	—
#5	$2n$	$2n$	—	0.3	—
#6	n^*	n	—	0.6	—
#7	$2n$	$2n$	—	0.4	—
#8	$2n$	$2n$	—	0.7	—

Table 1. Experimental evaluation on selected examples used for evaluation of Loopus [4]. Benchmarks are publicly available at [bitbucket](https://bitbucket.org).

at lines 22 and 27. If user function call appears in the analyzed code during analysis, like at line 26 of our example, the analyzer is provided with a summary of the function if available or the function is analyzed on demand. The summary is then applied to an abstract state at a call site. So in our example summary of `foo` will be applied to the abstract state of `thread1`.

Listing 2. Simple example capturing a deadlock between two global locks in C language using POSIX threads execution model

```

16 void foo() {
17     pthread_mutex_lock(&lock2);
18 }
21 void *thread1(...) {
22     pthread_mutex_lock(&lock1);
23     :
26     foo();
27     pthread_mutex_unlock(&lock1);
28 }
29 void *thread2(...) {
30     pthread_mutex_lock(&lock2);
31     :
36     pthread_mutex_lock(&lock1);
37 }
```

Next L2D2 looks through all the summaries of analyzed program and checks whether a potential deadlock can occur by computing transitive closure of relation consisting of all dependencies (see Listing 3) and checking if any lock depends on itself. The summaries for functions from the above example record information about the state of locks `lock1` and `lock2` as follows:

Listing 3. Summaries of functions in Listing 2

```

foo()
PRECONDITION: { unlocked={lock2} }
```

there are a few that meet the scalability condition, like *starvation* analyzer implemented in Facebook Infer. The problem of this analyzer is that it uses heuristic on the class root of the access path of the lock so it doesn't handle a pure C lock. Also worth mentioning is the RacerX analyzer [6], which is based on counting so called *locksets* i.e. sets of locks currently held. RacerX uses interprocedural, flow-sensitive and context-sensitive analysis. What means that each function needs to be reanalysed in a new context. Hence, we decide to adapt lockset analysis from RacerX to follow principles of Facebook Infer and by that create context-insensitive analysis which will be faster and more scalable. So we present Low Level Deadlock Detector (L2D2), the principle of which will be illustrate with the example in Listing 2.

L2D2 works by first computing a summary for each function by looking for lock and unlock events. Example of lock and unlock is illustrated in Listing 2

```

339 POSTCONDITION: { lockset={lock2} }
340 thread1(...)
341 PRECONDITION: { unlocked={lock1, lock2} }
342 POSTCONDITION: {
343     lockset={lock1, lock2},
344     dependencies={lock1->lock2}
345 }
346 thread2(...)
347 PRECONDITION: { unlocked={lock1, lock2} }
348 POSTCONDITION: {
349     lockset={lock1, lock2},
350     dependencies={lock2->lock1}
351 }

```

352 If we run L2D2 on code from our example it will
353 report a possible deadlock between two threads due
354 to cyclic dependency between $\text{lock1} \rightarrow \text{lock2}$ and
355 $\text{lock2} \rightarrow \text{lock1}$ that arises if thread 1 holds lock1
356 and waits on lock2 and thread 2 hold lock2 and
357 waits on lock1 .

358 4.1 Computing procedure summaries

359 In this subsection, we describe structure of the sum-
360 mary and process of computing it. To detect potential
361 deadlock we need to record information that will allow
362 us to answer these questions:

- 363 (1) What is the state of locks used in the analyzed
364 program?
- 365 (2) Could cyclic dependency on pending threads
366 occur?

367 To answer question (1), we have defined sets *lock-*
368 *set* and *unlockset*, which contains currently locked
369 and unlocked locks respectively. We have also added
370 sets *locked* and *unlocked* that serve as a precondition
371 for a given function and contain locks that should be
372 locked/unlocked before calling this function. Semantic
373 of these sets is as follows:

```

374 semantics of lockset:
375   lock(l) → lockset = lockset ∪ {l}
376   unlock(l) → lockset = lockset - {l}
377 semantics of unlockset:
378   lock(l) → unlockset = unlockset - {l}
379   unlock(l) → unlockset = unlockset ∪ {l}
380 semantics of locked:
381   if(lock(l) is first operation in f)
382     unlockedf = unlockedf ∪ {l}
383 semantics of unlocked:
384   if(unlock(l) is first operation in f)
385     lockedf = lockedf ∪ {l}

```

386 The summary also contains a set of one-level *de-*
387 *pendencies* by using which we can answer (2)nd ques-
388 tion. Extraction of these *dependencies* is called on
389 every lock acquisition and iterates over every lock in
390 the current *lockset*, emitting the ordering constraint
391 produced by the current acquisition. For example, if
392 lock2 is in the current *lockset* and lock1 has just

393 been acquired, the dependency $\text{lock2} \rightarrow \text{lock1}$ will
394 be emitted, as we can see in Listing 2 in function
395 thread2.

396 The most difficult part of dependencies extraction
397 is elimination of false ones caused by invalid *locksets*.
398 The main reasons for errors in *locksets* include the
399 number of conditionals, function calls and degree of
400 aliasing involved

401 *Applying procedure summaries.* As we mentioned
402 at the beginning of this section, if a function call ap-
403 pears in an analyzed code, we have to apply a summary
404 of the function to an abstract state at a callsite. Given
405 callee g , its lockset L_g , unlockset U_g and caller f , its
406 lockset L_f and unlockset U_f , we:

- 407 (1) Update the summary of g by replacing formal
408 parameters with actual ones in case that locks
409 were passed to the g as parameters. In the exam-
410 ple below, you can notice that in the summary
411 of g will be lock4 replaced with lock2 .
- 412 (2) Update the precondition of f :
413 $\text{if}(\exists l : l \in \text{unlocked}_g \wedge l \notin \text{unlockset}_f)$
414 add lock l to unlocked_f
415 $\text{if}(\exists l : l \in \text{locked}_g \wedge l \notin \text{lockset}_f)$
416 add lock l to locked_f
- 417 (3) Update the L_f : $L_f = (L_f \setminus U_g) \cup L_g$
- 418 (4) Update the U_f : $U_f = (U_f \setminus L_g) \cup U_g$
- 419 (5) Update the dependencies of f by adding new
420 dependencies for all locks in the caller's lockset
421 with locks which were locked in the callee. But
422 what if all the locks that were acquired in the
423 callee were also released there, as we can see in
424 the example below.

```

425 void f() {
426     pthread_mutex_lock(&lock2);
427     g(&lock2);
428 }
429 void g(pthread_mutex_t *lock4) {
430     pthread_mutex_lock(&lock3);
431     pthread_mutex_unlock(lock4);
432     pthread_mutex_lock(&lock1);
433     :
434     pthread_mutex_unlock(&lock1);
435     pthread_mutex_unlock(&lock3);
436 }

```

437 In that case, the callee's lockset will be empty
438 and we have no information about these locks.
439 So we had to add a new set to the summary
440 which semantics is similar to the semantics of
441 lockset except that unlock statement does not
442 remove a lock from it. In our example, this set
443 would contain lock3 and lock1 but there is
444 still one problem left. What if the lock from the
445 current lockset was unlocked in the callee before
446 we locked another lock there? Then we will emit

the wrong dependency `lock2→lock1`. In order to avoid this, we create `unlock→lock` type dependencies in summary, that can be used to safely determine the order of an operation in the callee. So this ensures that the only newly created correct dependency in our example will be `lock2→lock3`.

4.2 Reporting deadlocks

For deadlock detection, we use algorithm that iterates through all the summaries and computes the transitive closure of all dependencies. It records the cyclic lock dependency and displays the results to the user for inspection. Each deadlock is normally reported twice, at each trace starting point. So in our example in Listing 2, will be the deadlock reported for the first time in function `thread1` and for the second time in function `thread2`. TsF: I'm thinking whether this could be defined as some kind of lemma. E.g "Lemma 1. Given function f and its summary f_s and its transitive closure f_s^* , f contains deadlock iff $\exists a \in \text{something} : (a, a) \in f_s^*$. No proof. Fuck proofs.

4.3 Experimental evaluation

We performed our experiments by using 1002 concurrent C programs, that contain locks from the Debian GNU/Linux distribution. All benchmarks are available online at [gitlab](#). These programs were used for experimental evaluation of Daniel Kroening's static deadlock analyser [citace] implemented in the CPROVER framework.

This benchmark set consists of 11.4 MLOC. Of all the programs, 994 are assumed to be deadlock-free and 8 of them have proved deadlock. Our experiments were run on a CORE i7-7700HQ at 2.80 GHz running Ubuntu 18.04 with 64-bit binaries with comparison to the CPROVER experiments which were run on a Xeon X5667 at 3 GHz running Fedora 20 with 64-bit binaries. In case of CPROVER were memory and CPU time restricted to 24GB and 1800 seconds per benchmark.

Results. Our analyzer as same as CPROVER correctly report all 8 potential deadlocks in benchmarks with known issues. Comparison of results for deadlock-free programs you can see in Table 2.

	proved	alarms	t/o	m/o	errors
CPROVER	292	114	453	135	0
L2D2	810	104	0	0	80

Table 2. Results for the programs without deadlock (t/o – timed out, m/o – out of memory)

As you can see L2D2 reported false alarms for 104 deadlock-free benchmarks what is 10 less than

CPROVER. A much larger difference can be seen in cases where it was proved that there was no deadlock. The difference here is up to 518 examples in favor of our analyzer. In case of L2D2 you can see 80 compilation errors that were caused by syntax that Infer does not support. The biggest difference between our analyzer and CPROVER is runtime. While our analyzer needed approximately 2 hours to perform the experiments, CPROVER needed about 300 hours.

There is still space for improving our analysis by reduction of false alarms. The main reason for such alarms is false dependencies. Reasons for their existence we mentioned in subsection 4.1 (4th paragraph). So eliminating false positives consists of techniques to eliminate false dependencies.

5. Atomicity Violations Analyzer

In concurrent programs, there are often *atomicity requirements* for an execution of specific sequences of instructions. Violating these requirements may cause many kinds of problems, such as a unexpected behaviour, exceptions, segmentation faults or other failures. *Atomicity violations* are usually not verified by compilers, unlike syntactic or some sorts of semantic rules. Atomicity requirements, in most cases, are not event documented. It means that typically only programmers must take care of following these requirements. In general, it is very difficult to avoid errors in *atomicity-related programs*, especially in large projects, and even harder and time-consuming is then finding and fixing these errors.

In this section of this paper, there is described a proposal and an implementation of an *static analyzer for finding atomicity violations*.

5.1 Contracts for Concurrency

The proposal of a solution is based on the concept of *contracts for concurrency* described in [7]. These contracts allow to define *sequences of functions* that are required to be *executed atomically*. The proposed analyzer itself is able to produce mentioned contracts, and then verify whether the contracts are fulfilled.

In [7], a *basic contract* is formally defined as follows. Let $\Sigma_{\mathbb{M}}$ be a set of all function names of a software module. A *contract* is a set \mathbb{R} of *clauses* where each clause $\varrho \in \mathbb{R}$ is a regular expression over $\Sigma_{\mathbb{M}}$. A *contract violation* occurs if any of the sequences represented by the contract clauses is interleaved with an execution of functions from $\Sigma_{\mathbb{M}}$.

Consider an implementation of a function that replaces item a in an array by item b , as illustrates Listing 4. The contract for this specific scenario contains

542 clause ϱ_1 , which is defined and follows:

(ϱ_1) index_of set

543 Clause ϱ_1 specifies that the execution of `index_of`
544 followed by execution of `set` should be atomic. The
545 index of an item in an array is acquired, and then the
546 index is used to modify the array. Without atomicity,
547 a concurrent modification of the array may change
548 a position of the item. The acquired index then may
549 be invalid when `set` is executed.

Listing 4. Example of a contract violation

```
550 void replace(int *array, int a, int b) {  
551     int i = index_of(array, a);  
552     if (i >= 0) {  
553         set(array, i, b);  
554     }  
555 }  
556
```

557 In paper [7], there is described a proposal and an
558 implementation for a static validation which is based
559 on *grammars* and *parsing trees*. Within paper [7]
560 was implemented a stand-alone prototype tool¹ for
561 analysing applications, written in Java language, which
562 obtained promising experimental results. However, we
563 decided to propose and implement the analysis quite
564 different way, see 5.2. Moreover, we decided to imple-
565 ment this solution in the *Facebook Infer*, i.e., widely
566 used, active and a open source tool. Therefore the
567 analysis should be faster and more scalable thanks to
568 the way how the Facebook Infer works, as it was de-
569 scribed in section 2. The implementation is aimed for
570 programs written in C/C++ languages using *POSIX*
571 *Threads* (*Pthreads*) locks for a *synchronization of con-*
572 *current threads*. We are also focusing to reduce false
573 positive errors.

574 In the Facebook Infer, there is already implemented
575 an analysis called *Lock Consistency Violation*, see ².
576 That analysis finds atomicity violations for writes/reads
577 single variables that are required to be executed atom-
578 ically. Ours analysis is more general because it finds
579 *atomicity violations for sequences of functions* that
580 are required to be executed atomically, i.e., it checks
581 whether contracts for concurrency are fulfilled.

582 5.2 Two-Phase Analysis

583 The proposed solution is divided into two parts (*phases*
584 *of analysis*):

¹<https://github.com/trxsys/gluon>

²[https://fbinfer.com/docs/
checkers-bug-types.html#LOCK_CONSISTENCY_
VIOLATION](https://fbinfer.com/docs/checkers-bug-types.html#LOCK_CONSISTENCY_VIOLATION)

Phase 1 Detection of *sequences of functions* that should
be *executed atomically*. 585

Phase 2 Detection of *atomicity violations*. 587

[[TODO]] 588

589 5.3 Future Work

[[TODO]] 590

6. Conclusions 591

[Paper Summary] What was the paper about, then? 592
What the reader needs to remember about it? Lorem 593
ipsum dolor sit amet, consectetur adipiscing elit. Proin 594
vitae aliquet metus. Sed pharetra vehicula sem ut var- 595
ius. Aliquam molestie nulla et mauris suscipit, ut 596
commodo nunc mollis. 597

[Highlights of Results] Exact numbers. Remind 598
the reader that the paper matters. Lorem ipsum do- 599
lor sit amet, consectetur adipiscing elit. Sed tempus 600
fermentum ipsum at venenatis. Curabitur ultricies, 601
mauris eu ullamcorper mattis, ligula purus dapibus mi, 602
vel dapibus odio nulla et ex. Sed viverra cursus mattis. 603
Suspendisse ornare semper condimentum. Interdum et 604
malesuada fames ac ante ipsum. 605

[Paper Contributions] What is the original con- 606
tribution of this work? Two or three thoughts that one 607
should definitely take home. Lorem ipsum dolor sit 608
amet, consectetur adipiscing elit. Praesent posuere 609
mattis ante at imperdiet. Cras id tincidunt purus. Ali- 610
quam erat volutpat. Morbi non gravida nisi, non iaculis 611
tortor. Quisque at fringilla neque. 612

[Future Work] How can other researchers / devel- 613
opers make use of the results of this work? Do you 614
have further plans with this work? Or anybody else? 615
Lorem ipsum dolor sit amet, consectetur adipiscing 616
elit. Suspendisse sollicitudin posuere massa, non con- 617
vallis purus ultricies sit amet. Duis at nisl tincidunt, 618
maximus risus a, aliquet massa. Vestibulum libero 619
odio, condimentum ut ex non, eleifend. 620

Acknowledgements 621

We would like to thank our supervisors and colleagues 622
from VeriFIT team – Tomáš Fiedor, Adam Rogalewicz 623
and Tomáš Vojnar. We would also like to thank for 624
the support received from the H2020 ECSEL project 625
Aguas. 626

References 627

[1] Cristiano Calcagno, Dino Distefano, Peter 628
O'Hearn, and Hongseok Yang. *Compositional* 629
Shape Analysis by means of Bi-Abduction. In *Proc.* 630
of POPL'09, pages 289–300. 2009. 631

- 632 [2] Guoliang Jin, Linhai Song, Xiaoming Shi, Joel
633 Scherpelz, and Shan Lu. *Understanding and De-*
634 *tecting Real-World Performance Bugs. In Proc. of*
635 *PLDI'18*, pages 77–88. 2012.
- 636 [3] Stefan Bygde. *Static WCET analysis based on*
637 *abstract interpretation and counting of elements.*
638 PhD thesis, Mälardalen University, 2010.
- 639 [4] Moritz Sinn. *Automated Complexity Analysis for*
640 *Imperative Programs.* PhD thesis, Vienna Univer-
641 sity of Technology, 2016.
- 642 [5] A.H. Dogru and V. Bicar. *Modern Software En-*
643 *gineering Concepts and Practices: Advanced Ap-*
644 *proaches.* Premier reference source. Information
645 Science Reference, 2011.
- 646 [6] Dawson R. Engler and Ken Ashcraft. Rac-
647 erx: Effective, static detection of race conditions
648 and deadlocks. In *Operating Systems Review –*
649 *SIGOPS*, volume 37, pages 237–252, January
650 2003. DOI: 10.1145/1165389.945468.
- 651 [7] Ricardo J. Dias, Carla Ferreira, Jan Fiedor,
652 João M. Lourenço, Aleš Smrčka, Diogo G. Sousa,
653 and Tomáš Vojnar. Verifying concurrent pro-
654 grams using contracts. In *2017 IEEE Interna-*
655 *tional Conference on Software Testing, Verifica-*
656 *tion and Validation (ICST)*, pages 196–206. IEEE,
657 March 2017. DOI: 10.1109/ICST.2017.25. ISBN:
658 9781509060313.